

exercise__04_LauraHaegePhilippVonBachmann

June 3, 2021

1 Statistical Machine Learning Exercise Sheet 4

- Laura Haege
- Philipp Noel von Bachmann, Matrikelnummer: 4116220

1.1 Exercise 7

1.1.1 (a)

1. Suppose there exist y_1, y_2 such that $P_C(x) = \{y_1, y_2\}$ with $\|x - y_1\|_2^2 = \|x - y_2\|_2^2 = a$. Now construct $y^* = p \cdot y_1 + (1 - p) \cdot y_2 \in P_C(x)$

$$\begin{aligned} & \|x - p \cdot y_1 + (1 - p) \cdot y_2\|_2^2 \\ &= \|p \cdot x + (1 - p) \cdot x - (p \cdot y_1 + (1 - p) \cdot y_2)\|_2^2 \\ &= \|p \cdot x - p \cdot y_1 + (1 - p) \cdot x - (1 - p) \cdot y_2\|_2^2 \\ &\leq \|p \cdot x - p \cdot y_1\|_2^2 + \|(1 - p) \cdot x - (1 - p) \cdot y_2\|_2^2 \\ &= p^2 \cdot \|x - y_1\|_2^2 + (1 - p)^2 \cdot \|x - y_2\|_2^2 \\ &< p \cdot \|x - y_1\|_2^2 + (1 - p) \cdot \|x - y_2\|_2^2 \\ &= p \cdot a + (1 - p) \cdot a \\ &= a \end{aligned}$$

Therefore the distance to y^* is smaller than to y_1, y_2 and they can not be the minimum. Therefore the minimum has to be unique

2. As we measure the distance in L_2 , the distance within one component doesn't affect the distance within the other, and we can optimize each component separately. If $x_i < 0$, then the closest point in C is given by 0. Otherwise, we can just stay at x_i , which is the trivial minimal distance. This leads to $P_C(X)_i = \max(x_i, 0)$. In total we get $P_C(X) = \max(x, 0)$, where max is the element wise max.

2 (b)

- 1.

```
[1]: from sys import prefix
import numpy as np
import matplotlib.pyplot as plt
```

```

def LassoObjective(wplus, wminus, Phi, Y, lmbd):
    ''' evaluates the objective function at (wplus, wminus)
        L2 loss + L1 regularization
    '''
    w = wplus - wminus
    return ((Phi @ w - Y) ** 2).mean(
        ) + lmbd * np.abs(w).sum()

def GradLassoObjective(wplus, wminus, Phi, Y, lmbd):
    ''' computes the gradients of the objective function
        at (wplus, wminus)
        gradwplus: gradient wrt wplus
        gradwminus: gradient wrt minus

        FILL IN
    '''
    prediction_loss = 2* (Y - Phi @ wplus + Phi @ wminus)
    gradwplus = np.expand_dims(- np.sum(np.multiply(prediction_loss, Phi),
→axis=0) / Phi.shape[0] + lmbd, axis=1)
    gradwminus = np.expand_dims(np.sum(np.multiply(prediction_loss, Phi),
→axis=0) / Phi.shape[0] + lmbd, axis=1)
    return gradwplus, gradwminus

def ProjectionPositiveOrthant(x):
    ''' returns the projection of x onto the positive orthant

        FILL IN
    '''
    return np.maximum(0, x)

def getStepSize(wplus, wminus, Phi, Y, lmbd, gradwplus,
                gradwminus, loss):
    ''' performs one step of projected gradient descent (i.e.
        compute next iterate) with step size selection via
        backtracking line search

        input
        loss: objective function at current iterate (wplus, wminus)

        output
        wplusnew, wminusnew: next iterates wplus_{t+1}, wminus_{t+1}
        lossnew: objective function at the new iterate
    '''

```

```

FILL IN
'''
alpha, beta, sigma = 1., .1, .1
wplusnew, wminusnew = wplus.copy(), wminus.copy()
lossnew = np.float('Inf') # make sure to enter the loop

# choose the step size alpha with backtracking line search
while lossnew > loss + sigma * ((gradwplus * (
    wplusnew - wplus)).sum() + (gradwminus * (
    wminusnew - wminus)).sum()):
    # get new step size to test
    alpha *= beta

    # projected gradient step for wplus and wminus with step size alpha
    # i.e. compute  $x_{t+1}$  as in the text
    # FILL IN
    # print('fill in with projected gradient step')
    wplusnew = ProjectionPositiveOrthant(wplus - alpha * gradwplus)
    wminusnew = ProjectionPositiveOrthant(wminus - alpha * gradwminus)

    # compute new value of the objective
    lossnew = LassoObjective(wplusnew, wminusnew, Phi, Y, lmbd)

return wplusnew, wminusnew, lossnew

def Lasso(Phi, Y, lmbd):
    ''' compute weight of linear regression with Lasso

    Phi: design matrix  $n \times d$ 
    Y: true values  $n \times 1$ 
    lmbd: weight of regularization

    output: weights of linear regression  $d \times 1$ 
    '''
    # initialize wplus, wminus
    wplus = np.random.rand(Phi.shape[1], 1)
    wminus = np.random.rand(*wplus.shape)
    loss = LassoObjective(wplus, wminus, Phi, Y, lmbd)

    counter = 1
    while counter > 0:
        # compute gradients wrt wplus and wminus
        gradwplus, gradwminus = GradLassoObjective(
            wplus, wminus, Phi, Y, lmbd)

        # compute new iterates

```

```

wplus, wminus, loss = getStepSize(wplus,
    wminus, Phi, Y, lmbd, gradwplus, gradwminus, loss)

if (counter % 100) == 0:
    # check if stopping criterion is met
    wnew = wplus - wminus
    ind = wnew != 0.
    indz = wnew == 0.
    r = 2 / Phi.shape[0] * (Phi.T @ (Phi @ wnew - Y))
    stop = np.abs(r[ind] + lmbd * np.sign(wnew[ind]
        )).sum() + (np.abs(r[indz]) - lmbd * np.ones_like(
            r[indz])).clip(0.).sum()
    print('iter={} current objective={:.3f} nonzero weights={}'.format(
        counter, loss, ind.sum()) + \
        ' stop={:.5f}'.format(stop / Phi.shape[0]))
    if np.abs(stop) / Phi.shape[0] < 1e-5:
        break
    counter += 1

return wplus - wminus

```

2.

```

[2]: # load data
data = np.load("multidim_data_trainval.npy", allow_pickle=True).item()
x_train = data["Xtrain"]
y_train = data["Ytrain"]
x_val = data["Xval"]
y_val = data["Yval"]

# normalize
x_train -= np.mean(x_train, axis=0)
x_train /= np.std(x_train, axis=0)
x_val -= np.mean(x_val, axis=0)
x_val /= np.std(x_val, axis=0)

# add one vector for offset
x_train = np.hstack((np.ones((x_train.shape[0], 1)), x_train))
x_val = np.hstack((np.ones((x_val.shape[0], 1)), x_val))

```

```

[3]: lmbd = 10

w = Lasso(x_train, y_train, lmbd=lmbd)
y_train_predict = x_train @ w
y_val_predict = x_val @ w
train_loss = np.linalg.norm(y_train - y_train_predict)**2/x_train.shape[0]

```

```
test_loss = np.linalg.norm(y_val - y_val_predict)**2/x_val.shape[0]
```

```
iter=100 current objective=139277.680 nonzero weights=63 stop=0.11879
iter=200 current objective=138925.125 nonzero weights=45 stop=0.03480
iter=300 current objective=138832.709 nonzero weights=44 stop=0.00569
iter=400 current objective=138786.650 nonzero weights=43 stop=0.00400
iter=500 current objective=138758.119 nonzero weights=43 stop=0.00305
iter=600 current objective=138738.583 nonzero weights=42 stop=0.00240
iter=700 current objective=138724.429 nonzero weights=42 stop=0.00194
iter=800 current objective=138716.244 nonzero weights=41 stop=0.00125
iter=900 current objective=138714.441 nonzero weights=41 stop=0.00078
iter=1000 current objective=138713.814 nonzero weights=40 stop=0.00048
iter=1100 current objective=138713.589 nonzero weights=40 stop=0.00030
iter=1200 current objective=138713.506 nonzero weights=40 stop=0.00018
iter=1300 current objective=138713.476 nonzero weights=40 stop=0.00011
iter=1400 current objective=138713.464 nonzero weights=40 stop=0.00007
iter=1500 current objective=138713.460 nonzero weights=40 stop=0.00004
iter=1600 current objective=138713.458 nonzero weights=40 stop=0.00003
iter=1700 current objective=138713.458 nonzero weights=40 stop=0.00002
iter=1800 current objective=138713.457 nonzero weights=40 stop=0.00001
iter=1900 current objective=138713.457 nonzero weights=40 stop=0.00001
```

i.


```
[4]: print(f"The training loss is {train_loss} and the test loss is {test_loss}")
```

The training loss is 121332.09617441503 and the test loss is 141717.19168760636

ii.

```
[5]: perfect_fit = np.linspace(0, 5000)
plt.scatter(y_val, y_val_predict)
plt.plot(perfect_fit, perfect_fit, color="red")
```

```
[5]: [<matplotlib.lines.Line2D at 0x1d693d40708>]
```



We can see that for small values, the predictions are very good. However for larger values, the predictions get more inaccurate

```
[6]: def basis(X):
      return np.hstack([X, np.sin(np.pi * X), np.cos(np.pi * X)])
```

```
[7]: def train(X, Y):
      X = basis(X)
      w = Lasso(X, Y, lmbd=5)
      return w
```

```
[8]: w_predict = train(x_train, y_train)
```

```
iter=100 current objective=121815.273 nonzero weights=229 stop=0.28289
iter=200 current objective=120535.502 nonzero weights=191 stop=0.20205
iter=300 current objective=120065.688 nonzero weights=195 stop=0.12329
iter=400 current objective=119775.663 nonzero weights=184 stop=0.13487
iter=500 current objective=119604.110 nonzero weights=185 stop=0.06124
iter=600 current objective=119506.657 nonzero weights=180 stop=0.04667
iter=700 current objective=119453.600 nonzero weights=180 stop=0.02954
iter=800 current objective=119413.800 nonzero weights=177 stop=0.04661
iter=900 current objective=119381.764 nonzero weights=179 stop=0.03385
iter=1000 current objective=119354.426 nonzero weights=179 stop=0.03198
iter=1100 current objective=119330.297 nonzero weights=179 stop=0.02725
iter=1200 current objective=119308.427 nonzero weights=179 stop=0.02730
iter=1300 current objective=119288.177 nonzero weights=176 stop=0.03379
```

iter=1400	current	objective=119269.033	nonzero	weights=177	stop=0.01911
iter=1500	current	objective=119250.651	nonzero	weights=177	stop=0.03064
iter=1600	current	objective=119232.918	nonzero	weights=179	stop=0.02207
iter=1700	current	objective=119215.704	nonzero	weights=177	stop=0.01735
iter=1800	current	objective=119198.842	nonzero	weights=178	stop=0.03668
iter=1900	current	objective=119182.199	nonzero	weights=176	stop=0.03053
iter=2000	current	objective=119165.837	nonzero	weights=177	stop=0.02698
iter=2100	current	objective=119149.694	nonzero	weights=180	stop=0.02393
iter=2200	current	objective=119137.323	nonzero	weights=176	stop=0.01101
iter=2300	current	objective=119130.836	nonzero	weights=176	stop=0.01331
iter=2400	current	objective=119124.776	nonzero	weights=176	stop=0.01838
iter=2500	current	objective=119118.961	nonzero	weights=177	stop=0.01049
iter=2600	current	objective=119113.316	nonzero	weights=175	stop=0.01290
iter=2700	current	objective=119107.859	nonzero	weights=175	stop=0.01615
iter=2800	current	objective=119102.526	nonzero	weights=175	stop=0.01995
iter=2900	current	objective=119097.244	nonzero	weights=175	stop=0.00973
iter=3000	current	objective=119092.038	nonzero	weights=176	stop=0.01127
iter=3100	current	objective=119086.903	nonzero	weights=175	stop=0.01256
iter=3200	current	objective=119081.836	nonzero	weights=178	stop=0.01422
iter=3300	current	objective=119076.828	nonzero	weights=174	stop=0.01559
iter=3400	current	objective=119071.889	nonzero	weights=174	stop=0.01704
iter=3500	current	objective=119067.196	nonzero	weights=173	stop=0.01763
iter=3600	current	objective=119062.598	nonzero	weights=174	stop=0.00761
iter=3700	current	objective=119058.036	nonzero	weights=174	stop=0.00804
iter=3800	current	objective=119053.512	nonzero	weights=174	stop=0.00838
iter=3900	current	objective=119049.018	nonzero	weights=174	stop=0.00883
iter=4000	current	objective=119044.552	nonzero	weights=174	stop=0.00928
iter=4100	current	objective=119040.110	nonzero	weights=174	stop=0.00975
iter=4200	current	objective=119035.691	nonzero	weights=174	stop=0.01032
iter=4300	current	objective=119032.790	nonzero	weights=172	stop=0.00415
iter=4400	current	objective=119032.324	nonzero	weights=172	stop=0.00458
iter=4500	current	objective=119032.011	nonzero	weights=172	stop=0.00219
iter=4600	current	objective=119031.769	nonzero	weights=172	stop=0.00227
iter=4700	current	objective=119031.567	nonzero	weights=172	stop=0.00245
iter=4800	current	objective=119031.392	nonzero	weights=171	stop=0.00251
iter=4900	current	objective=119031.235	nonzero	weights=171	stop=0.00263
iter=5000	current	objective=119031.094	nonzero	weights=171	stop=0.00276
iter=5100	current	objective=119030.964	nonzero	weights=171	stop=0.00123
iter=5200	current	objective=119030.845	nonzero	weights=171	stop=0.00128
iter=5300	current	objective=119030.735	nonzero	weights=171	stop=0.00134
iter=5400	current	objective=119030.633	nonzero	weights=170	stop=0.00140
iter=5500	current	objective=119030.538	nonzero	weights=170	stop=0.00147
iter=5600	current	objective=119030.451	nonzero	weights=170	stop=0.00154
iter=5700	current	objective=119030.369	nonzero	weights=170	stop=0.00162
iter=5800	current	objective=119030.294	nonzero	weights=170	stop=0.00170
iter=5900	current	objective=119030.224	nonzero	weights=170	stop=0.00180
iter=6000	current	objective=119030.159	nonzero	weights=170	stop=0.00189
iter=6100	current	objective=119030.099	nonzero	weights=170	stop=0.00083

```

iter=6200 current objective=119030.043 nonzero weights=170 stop=0.00088
iter=6300 current objective=119029.990 nonzero weights=170 stop=0.00091
iter=6400 current objective=119029.942 nonzero weights=170 stop=0.00096
iter=6500 current objective=119029.919 nonzero weights=169 stop=0.00069
iter=6600 current objective=119029.908 nonzero weights=169 stop=0.00043
iter=6700 current objective=119029.900 nonzero weights=169 stop=0.00054
iter=6800 current objective=119029.895 nonzero weights=169 stop=0.00032
iter=6900 current objective=119029.891 nonzero weights=169 stop=0.00020
iter=7000 current objective=119029.889 nonzero weights=169 stop=0.00028
iter=7100 current objective=119029.887 nonzero weights=169 stop=0.00017
iter=7200 current objective=119029.886 nonzero weights=169 stop=0.00023
iter=7300 current objective=119029.885 nonzero weights=169 stop=0.00014
iter=7400 current objective=119029.885 nonzero weights=169 stop=0.00009
iter=7500 current objective=119029.884 nonzero weights=169 stop=0.00012
iter=7600 current objective=119029.884 nonzero weights=169 stop=0.00007
iter=7700 current objective=119029.884 nonzero weights=169 stop=0.00010
iter=7800 current objective=119029.884 nonzero weights=169 stop=0.00006
iter=7900 current objective=119029.883 nonzero weights=169 stop=0.00008
iter=8000 current objective=119029.883 nonzero weights=169 stop=0.00005
iter=8100 current objective=119029.883 nonzero weights=169 stop=0.00003
iter=8200 current objective=119029.883 nonzero weights=169 stop=0.00004
iter=8300 current objective=119029.883 nonzero weights=169 stop=0.00003
iter=8400 current objective=119029.883 nonzero weights=169 stop=0.00003
iter=8500 current objective=119029.883 nonzero weights=169 stop=0.00002
iter=8600 current objective=119029.883 nonzero weights=169 stop=0.00003
iter=8700 current objective=119029.883 nonzero weights=169 stop=0.00002
iter=8800 current objective=119029.883 nonzero weights=169 stop=0.00001
iter=8900 current objective=119029.883 nonzero weights=169 stop=0.00001
iter=9000 current objective=119029.883 nonzero weights=169 stop=0.00001

```

```

[9]: def Prediction(X):
      return basis(X) @ w_predict

```

2.1 Exercise 8

2.1.1 (a)

The minimization problem is convex because the primal problem is convex (L_2 norm is convex), and the equality constrain $Xw = b$ is linear.

2.1.2 (b)

First step: Show that the solution satisfies the constrain:

$$\begin{aligned} X(w_0 - X^T(XX^T)^{-1}(Xw_0 - b)) &= Xw_0 - X(X^T(XX^T)^{-1}(Xw_0 - b)) \\ &= Xw_0 - XX^T(XX^T)^{-1}(Xw_0 - b) \\ &= Xw_0 - (Xw_0 - b) \\ &= Xw_0 - Xw_0 + b \\ &= b \end{aligned}$$

Now setup Lagrange Dual Problem:

$$L(w, v) = \|w - w_0\| + v^T(Xw - b)$$

Compute derivative

$$\begin{aligned} \nabla_w L(w, v) &= \nabla_w \|w - w_0\| + v^T(Xw - b) \\ &= 2(w - w_0) + v^T X \end{aligned}$$

Set to 0:

$$\begin{aligned} 0 &= 2(w - w_0) + v^T X \\ w &= -v^T X \frac{1}{2} + w_0 \end{aligned}$$

Plug in L to get g .

$$\begin{aligned} g(v) &= L(-v^T X \frac{1}{2} + w_0, v) = \|-v^T X \frac{1}{2} + w_0 - w_0\| + v^T(X(-v^T X \frac{1}{2} + w_0) - b) \\ &= \frac{1}{4}\|v^T X\| + v^T(X(-v^T X \frac{1}{2} + w_0) - b) \end{aligned}$$